

# loantap

September 30, 2024

## 0.1 About LoanTap

LoanTap is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer-friendly terms to salaried professionals and businessmen.

### 0.1.1 Business Problem

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals.

LoanTap deploys formal credit to salaried individuals and businesses through 4 main financial instruments:

- Personal Loan
- EMI Free Loan
- Personal Overdraft
- Advance Salary Loan

This case study will focus on the underwriting process behind Personal Loan only.

### 0.1.2 Dataset

Column Name	Description
loan_amnt	The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.
term	The number of payments on the loan. Values are in months and can be either 36 or 60.
int_rate	Interest Rate on the loan.
installment	The monthly payment owed by the borrower if the loan originates.
grade	LoanTap assigned loan grade.
sub_grade	LoanTap assigned loan subgrade.
emp_title	The job title supplied by the Borrower when applying for the loan.
emp_length	Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.
home_ownership	The home ownership status provided by the borrower during registration or obtained from the credit report.
annual_inc	The self-reported annual income provided by the borrower during registration.

Column Name	Description
verification_status	Indicates if income was verified by LoanTap, not verified, or if the income source was verified.
issue_d	The month which the loan was funded.
loan_status	Current status of the loan - Target Variable.
purpose	A category provided by the borrower for the loan request.
title	The loan title provided by the borrower.
dti	A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LoanTap loan, divided by the borrower's self-reported monthly income.
earliest_cr_line	The month the borrower's earliest reported credit line was opened.
open_acc	The number of open credit lines in the borrower's credit file.
pub_rec	Number of derogatory public records.
revol_bal	Total credit revolving balance.
revol_util	Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolving credit.
total_acc	The total number of credit lines currently in the borrower's credit file.
initial_list_status	The initial listing status of the loan. Possible values are – W, F.
application_type	Indicates whether the loan is an individual application or a joint application with two co-borrowers.
mort_acc	Number of mortgage accounts.
pub_rec_bankruptcies	Number of public record bankruptcies.
Address	Address of the individual.

## Importing Required Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

import warnings
warnings.filterwarnings('ignore')

[2]: pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)

palette = ['#963f8d', '#624496']
```

```
sns.set(style='ticks', palette=palette)
```

## Read Dataset

```
[5]: df = pd.read_csv('../data/LoanTapData.csv')
df.sample(5)
```

```
[5]:      loan_amnt      term  int_rate  installment  grade  sub_grade  \
263539    20500.0   60 months    20.99         554.48     E         E4
25692     10000.0   36 months    10.99         327.34     B         B4
18277      6500.0   36 months    10.15         210.20     B         B2
327083     5350.0   36 months    12.12         178.01     B         B3
306382     1500.0   36 months     7.88          46.93     A         A5

      emp_title  emp_length  home_ownership  \
263539      Behavior Specialist    10+ years    MORTGAGE
25692              Rca    10+ years    MORTGAGE
18277      Payroll Admin     7 years      RENT
327083  peacock interiors and gallery inc     4 years      RENT
306382      Lawrence public schools     9 years      RENT

      annual_inc  verification_status  issue_d  loan_status  \
263539     57000.0      Source Verified  May-2014  Charged Off
25692     40000.0      Source Verified  Jul-2015   Fully Paid
18277     46000.0      Source Verified  Jul-2014   Fully Paid
327083     40000.0      Not Verified  Dec-2012   Fully Paid
306382     35000.0      Not Verified  Apr-2010   Fully Paid

      purpose      title  dti  earliest_cr_line  \
263539  debt_consolidation  Debt consolidation    9.41      Sep-1990
25692  debt_consolidation  Debt consolidation   15.30      May-2001
18277   credit_card  Credit card refinancing   32.92     Oct-1995
327083   credit_card  debit consolidation   13.50     Dec-2004
306382   vacation      kitten   11.69     Dec-1988

      open_acc  pub_rec  revol_bal  revol_util  total_acc  \
263539      9.0      5.0    4766.0      54.2      15.0
25692     13.0      2.0   10602.0      35.6      27.0
18277     13.0      0.0   14818.0      42.5      30.0
327083      7.0      0.0   16607.0      40.3      15.0
306382      9.0      0.0    3162.0      22.4      27.0

      initial_list_status  application_type  mort_acc  pub_rec_bankruptcies  \
263539                  w      INDIVIDUAL      1.0              0.0
25692                  w      INDIVIDUAL      5.0              2.0
18277                  f      INDIVIDUAL      0.0              0.0
327083                  f      INDIVIDUAL      0.0              0.0
306382                  f      INDIVIDUAL      NaN              0.0
```

```

                                address
263539      01246 Carrie Passage\r\nNew Kyle, ND 11650
25692    00296 Kirk Bypass Suite 650\r\nPerezburgh, LA ...
18277      81136 Beth Hollow\r\nRhondaland, RI 00813
327083    2423 Hines Spring Apt. 664\r\nNew Pamela, IL 0...
306382      302 Schmidt Avenue\r\nJeffreyville, NM 00813

```

```

[6]: print("Shape of the data: ", df.shape)
      print("The Given Dataset has {} rows and {} columns".format(df.shape[0], df.
        ↪shape[1]))
      print("Columns: ", df.columns.to_list())

```

Shape of the data: (396030, 27)

The Given Dataset has 396030 rows and 27 columns

Columns: ['loan\_amnt', 'term', 'int\_rate', 'installment', 'grade', 'sub\_grade', 'emp\_title', 'emp\_length', 'home\_ownership', 'annual\_inc', 'verification\_status', 'issue\_d', 'loan\_status', 'purpose', 'title', 'dti', 'earliest\_cr\_line', 'open\_acc', 'pub\_rec', 'revol\_bal', 'revol\_util', 'total\_acc', 'initial\_list\_status', 'application\_type', 'mort\_acc', 'pub\_rec\_bankruptcies', 'address']

### 0.1.3 Shape

- The dataset comprises 396030 rows and 27 columns, representing a volume of data.
- Each row corresponds to each loan distribution by the company.

### 0.1.4 Data Structure

```

[7]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             396030 non-null float64
1   term                  396030 non-null object
2   int_rate              396030 non-null float64
3   installment           396030 non-null float64
4   grade                 396030 non-null object
5   sub_grade             396030 non-null object
6   emp_title             373103 non-null object
7   emp_length           377729 non-null object
8   home_ownership        396030 non-null object
9   annual_inc            396030 non-null float64
10  verification_status   396030 non-null object
11  issue_d               396030 non-null object
12  loan_status           396030 non-null object

```

13	purpose	396030	non-null	object
14	title	394274	non-null	object
15	dti	396030	non-null	float64
16	earliest_cr_line	396030	non-null	object
17	open_acc	396030	non-null	float64
18	pub_rec	396030	non-null	float64
19	revol_bal	396030	non-null	float64
20	revol_util	395754	non-null	float64
21	total_acc	396030	non-null	float64
22	initial_list_status	396030	non-null	object
23	application_type	396030	non-null	object
24	mort_acc	358235	non-null	float64
25	pub_rec_bankruptcies	395495	non-null	float64
26	address	396030	non-null	object

dtypes: float64(12), object(15)  
memory usage: 81.6+ MB

```
[8]: df.isnull().sum()
```

```
[8]: loan_amnt      0
term              0
int_rate          0
installment       0
grade             0
sub_grade         0
emp_title        22927
emp_length       18301
home_ownership    0
annual_inc        0
verification_status  0
issue_d           0
loan_status       0
purpose           0
title            1756
dti              0
earliest_cr_line  0
open_acc         0
pub_rec          0
revol_bal        0
revol_util       276
total_acc        0
initial_list_status  0
application_type  0
mort_acc         37795
pub_rec_bankruptcies  535
address          0
dtype: int64
```

```
[9]: df_missing = df.isnull().sum()
df_missing = df_missing[df_missing > 0]
df_missing = df_missing.sort_values(ascending=False)
df_missing = df_missing.to_frame()
df_missing.columns = ['count']
df_missing
```

```
[9]:
```

	count
mort_acc	37795
emp_title	22927
emp_length	18301
title	1756
pub_rec_bankruptcies	535
revol_util	276

```
[10]: numeric_features = ['mort_acc', 'pub_rec_bankruptcies', 'revol_util']
df[numeric_features] = df[numeric_features].fillna(df[numeric_features].
↳median())

categorical_features = ['emp_title', 'emp_length', 'title']
df[categorical_features] = df[categorical_features].
↳fillna(df[categorical_features].mode().iloc[0])

df.isnull().sum()
```

```
[10]:
```

loan_amnt	0
term	0
int_rate	0
installment	0
grade	0
sub_grade	0
emp_title	0
emp_length	0
home_ownership	0
annual_inc	0
verification_status	0
issue_d	0
loan_status	0
purpose	0
title	0
dti	0
earliest_cr_line	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	0
total_acc	0

```

initial_list_status      0
application_type         0
mort_acc                0
pub_rec_bankruptcies    0
address                  0
dtype: int64

```

```
[11]: df.duplicated().sum()
```

```
[11]: np.int64(0)
```

```
[12]: df.describe().T
```

```
[12]:
```

	count	mean	std	min	25%	\
loan_amnt	396030.0	14113.888089	8357.441341	500.00	8000.00	
int_rate	396030.0	13.639400	4.472157	5.32	10.49	
installment	396030.0	431.849698	250.727790	16.08	250.33	
annual_inc	396030.0	74203.175798	61637.621158	0.00	45000.00	
dti	396030.0	17.379514	18.019092	0.00	11.28	
open_acc	396030.0	11.311153	5.137649	0.00	8.00	
pub_rec	396030.0	0.178191	0.530671	0.00	0.00	
revol_bal	396030.0	15844.539853	20591.836109	0.00	6025.00	
revol_util	396030.0	53.792451	24.443685	0.00	35.90	
total_acc	396030.0	25.414744	11.886991	2.00	17.00	
mort_acc	396030.0	1.736308	2.056819	0.00	0.00	
pub_rec_bankruptcies	396030.0	0.121483	0.355962	0.00	0.00	

	50%	75%	max
loan_amnt	12000.00	20000.00	40000.00
int_rate	13.33	16.49	30.99
installment	375.43	567.30	1533.81
annual_inc	64000.00	90000.00	8706582.00
dti	16.91	22.98	9999.00
open_acc	10.00	14.00	90.00
pub_rec	0.00	0.00	86.00
revol_bal	11181.00	19620.00	1743266.00
revol_util	54.80	72.90	892.30
total_acc	24.00	32.00	151.00
mort_acc	1.00	3.00	34.00
pub_rec_bankruptcies	0.00	0.00	8.00

```
[13]: df.describe(include='object').T
```

```
[13]:
```

	count	unique	top	freq
term	396030	2	36 months	302005
grade	396030	7	B	116018
sub_grade	396030	35	B3	26655

emp_title	396030	173105	Teacher	27316
emp_length	396030	11	10+ years	144342
home_ownership	396030	6	MORTGAGE	198348
verification_status	396030	3	Verified	139563
issue_d	396030	115	Oct-2014	14846
loan_status	396030	2	Fully Paid	318357
purpose	396030	14	debt_consolidation	234507
title	396030	48816	Debt consolidation	154228
earliest_cr_line	396030	684	Oct-2000	3017
initial_list_status	396030	2	f	238066
application_type	396030	3	INDIVIDUAL	395319
address	396030	393700	USS Smith\r\nFPO AP 70466	8

### 0.1.5 Dataset Information:

- **Data Consistency:** Few columns has missing values in the dataset.
- **Data Types:** Columns are classified into object and float types.
- **Duplicate:** There is no duplicate rows identified

### 0.1.6 Preprocessing

```
[14]: df['pincode'] = df['address'].apply(lambda x: x.split()[-1])
df['pincode'] = df['pincode'].astype('int')
```

```
[15]: df['state'] = df['address'].apply(lambda x: x.split()[-2])
df['state'] = df['state'].astype('str')
```

```
[16]: # Create flags
df['pub_rec_flag'] = df['pub_rec'].apply(lambda x: 1 if x > 1.0 else 0)
df['mort_acc_flag'] = df['mort_acc'].apply(lambda x: 1 if x > 1.0 else 0)
df['pub_rec_bankruptcies_flag'] = df['pub_rec_bankruptcies'].apply(lambda x: 1
↪ if x > 1.0 else 0)
```

```
[17]: df_ = df[df['loan_status']=='Fully Paid']
df_['emp_title'] = df_['emp_title'].str.lower()
(df_['emp_title'].value_counts(normalize=True) * 100).to_frame().head(10)
```

```
[17]: proportion
emp_title
teacher      6.702852
manager      1.356967
registered nurse  0.658066
supervisor   0.618174
sales         0.555980
rn            0.517344
driver        0.506350
owner         0.465201
project manager 0.463002
```



office manager            0.404891

Name the top 2 afforded job titles. - Teacher - Manager

```
[18]: # Columns Might not be useful
useless_columns = ['emp_title', 'issue_d', 'title', 'earliest_cr_line',
                  ↪ 'address', 'pub_rec', 'mort_acc', 'pub_rec_bankruptcies']
df = df.drop(useless_columns, axis=1)
```

```
[19]: # Object columns to Category
for col in df.select_dtypes(include='object').columns:
    df[col] = df[col].astype('category')
```

### 0.1.7 Insight

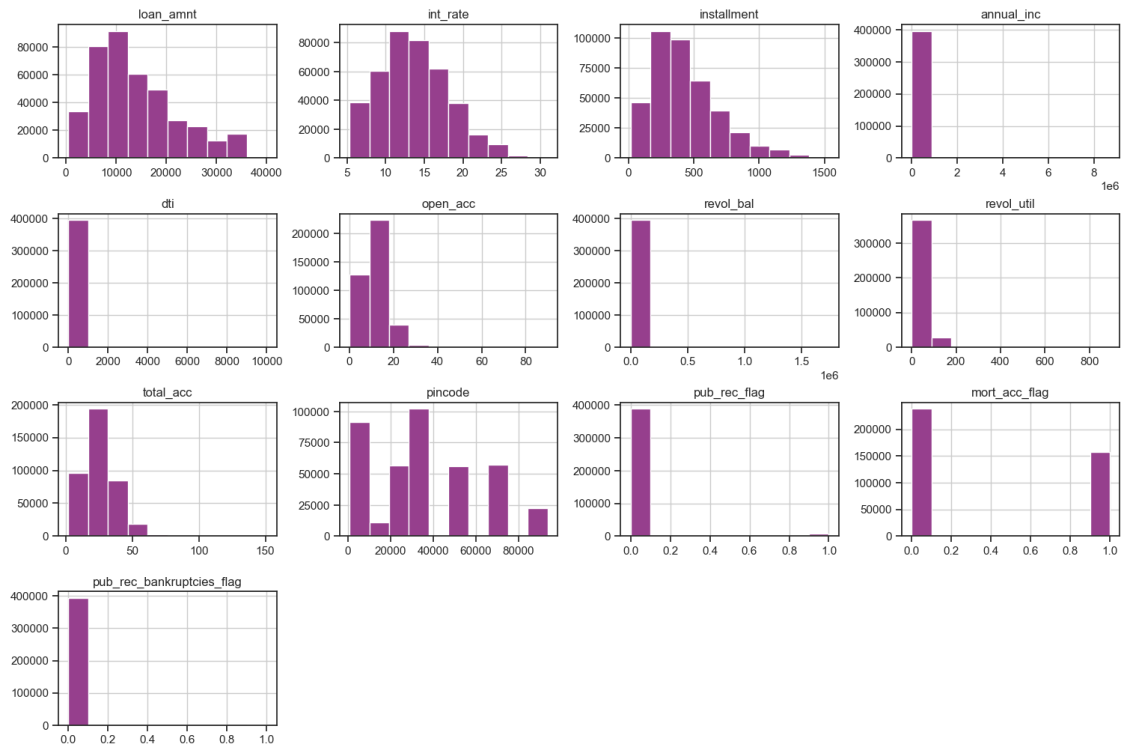
High cardinality columns are removed \* emp\_title (173105) \* title (48816) \* address (393700)

Other columns like \* issue\_d \* earliest\_cr\_line


and likely not useful for predicting loan status.

### 0.1.8 Exploratory Data Analysis (EDA)

```
[20]: df.hist(figsize=(15, 10), color=palette[0])
plt.subplots_adjust(hspace=0.5, wspace=0.5)
plt.xticks(fontsize=8)
plt.yticks(fontsize=8)
plt.tight_layout(pad=0.5)
plt.text(0.5, 0.5, 'Histograms', horizontalalignment='center',
        ↪ verticalalignment='center', fontsize=15, color='red', alpha=0.5)
plt.show()
```



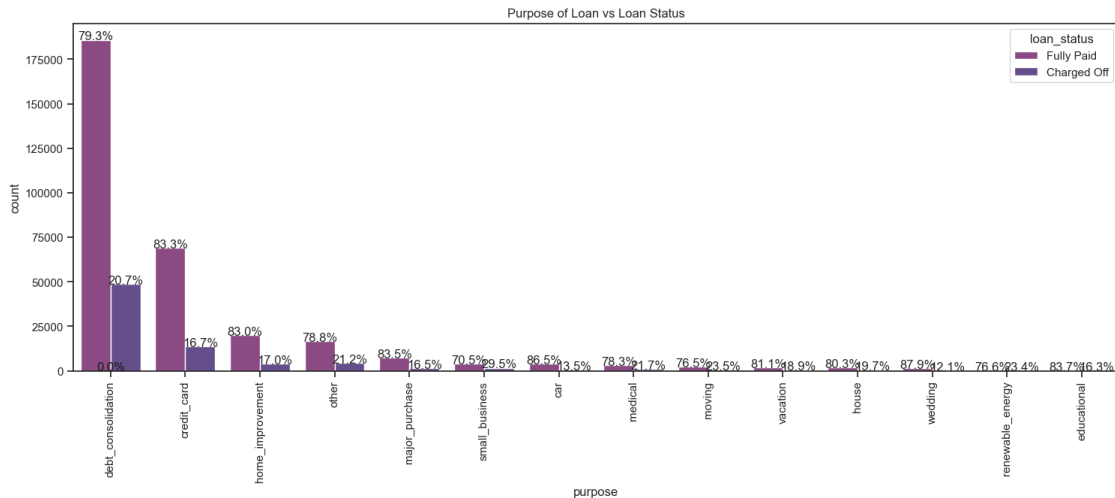
```
[21]: plt.figure(figsize=(10, 6))
sns.boxplot(data=df)
plt.xticks(rotation=45)
plt.show()
```



```
[22]: plt.figure(figsize=(18, 6))
loan_status_counts = df['loan_status'].value_counts()
sns.countplot(data=df, x='purpose', hue='loan_status',
              order=df['purpose'].value_counts().index,
              hue_order=loan_status_counts.index)
plt.xticks(rotation=90)
plt.title('Purpose of Loan vs Loan Status')

# Calculate percentages
total_counts = df['purpose'].value_counts()
for p in plt.gca().patches:
    height = p.get_height()
    total = total_counts[round(p.get_x())]
    percentage = height / total * 100
    plt.gca().text(p.get_x() + p.get_width() / 2, height + 20, f'{percentage:.1f}%', ha='center')

plt.show()
```



```
[23]: loan_status_by_purpose = df[['loan_status', 'purpose']].groupby(['purpose', 'loan_status']).size().unstack().reset_index()
# loan_status_by_purpose.drop(columns=['loan_status'], inplace=True)
loan_status_by_purpose['Total'] = loan_status_by_purpose[['Charged Off', 'Fully Paid']].sum(axis=1)
loan_status_by_purpose['Charged Off Percentage'] = round(loan_status_by_purpose['Charged Off'] / loan_status_by_purpose['Total'] * 100, 2)
loan_status_by_purpose['Fully Paid Percentage'] = round(loan_status_by_purpose['Fully Paid'] / loan_status_by_purpose['Total'] * 100, 2)
loan_status_by_purpose.sort_values(by='Charged Off Percentage', ascending=False, inplace=True)
loan_status_by_purpose
```

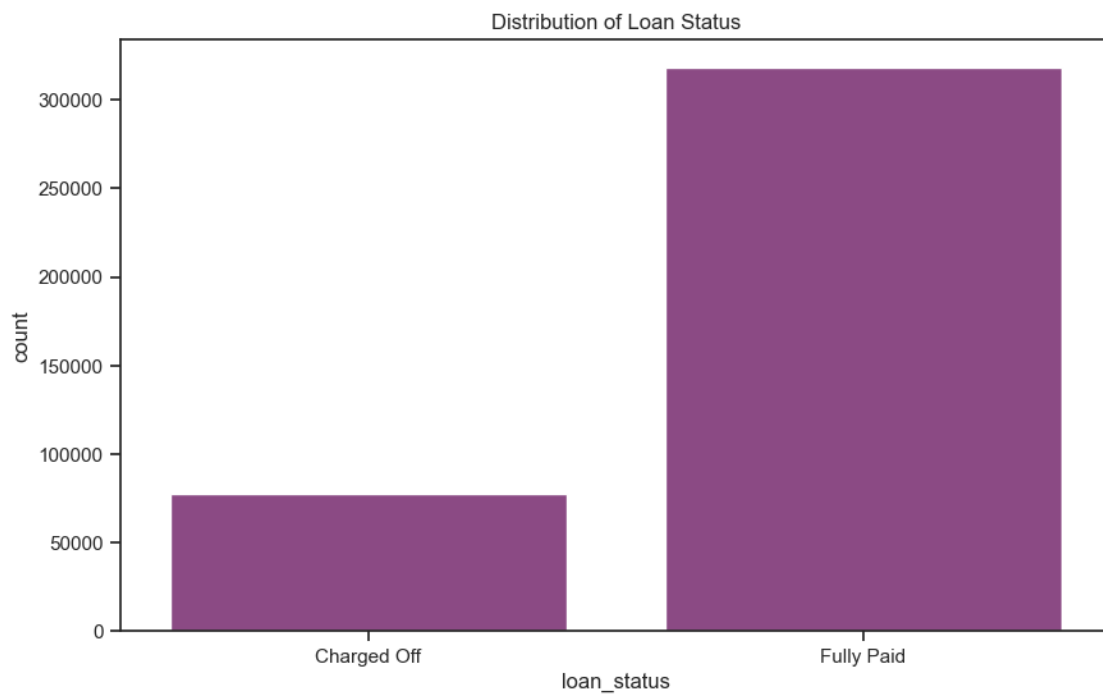
```
[23]: loan_status      purpose  Charged Off  Fully Paid  Total  \
11      small_business      1679      4022      5701
8          moving          670      2184      2854
10     renewable_energy          77       252       329
7          medical          911      3285      4196
9          other          4495     16690     21185
2     debt_consolidation     48640    185867    234507
5          house          434      1767       2201
12         vacation          464      1988       2452
4     home_improvement      4087     19943     24030
1          credit_card     13874     69145     83019
6     major_purchase      1448       7342       8790
3          educational          42        215        257
0          car           633      4064      4697
```

13	wedding	219	1593	1812
----	---------	-----	------	------

loan_status	Charged Off Percentage	Fully Paid Percentage
11	29.45	70.55
8	23.48	76.52
10	23.40	76.60
7	21.71	78.29
9	21.22	78.78
2	20.74	79.26
5	19.72	80.28
12	18.92	81.08
4	17.01	82.99
1	16.71	83.29
6	16.47	83.53
3	16.34	83.66
0	13.48	86.52
13	12.09	87.91

```
[24]: plt.figure(figsize=(10, 6))
sns.countplot(x='loan_status', data=df)
plt.title('Distribution of Loan Status')
plt.show()
```

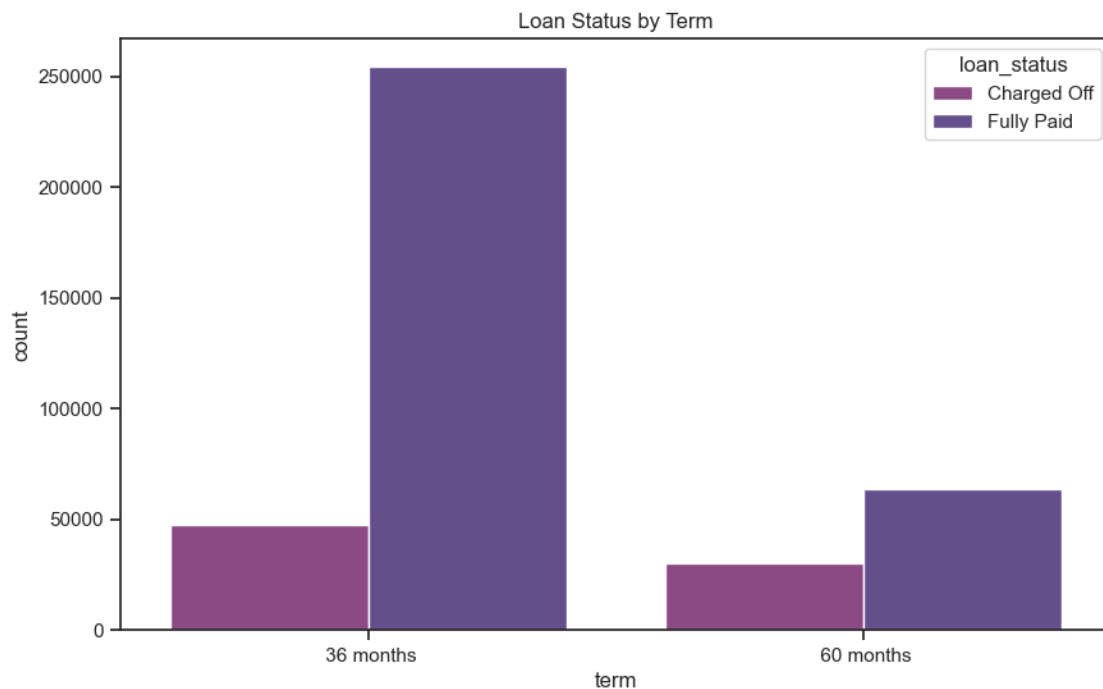


```
[25]: df['loan_status'].value_counts(normalize=True) * 100
```

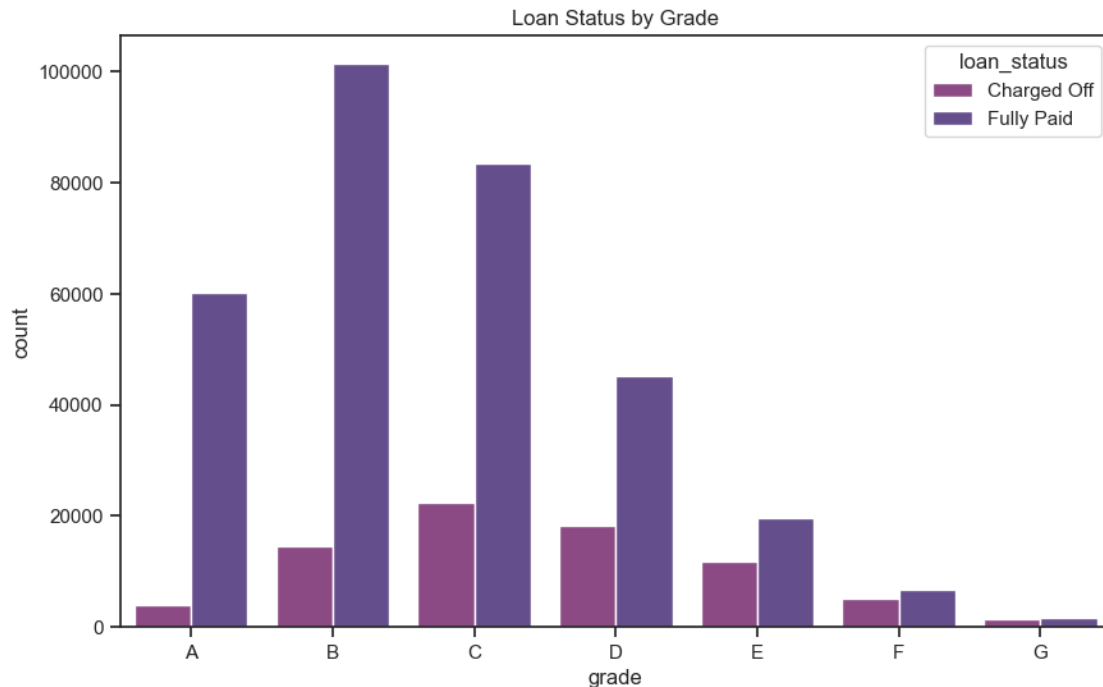
```
[25]: loan_status
Fully Paid      80.387092
Charged Off     19.612908
Name: proportion, dtype: float64
```

**What percentage of customers have fully paid their Loan Amount?** - 80.38% of loans are fully paid

```
[26]: plt.figure(figsize=(10, 6))
sns.countplot(x='term', hue='loan_status', data=df)
plt.title('Loan Status by Term')
plt.show()
```



```
[27]: plt.figure(figsize=(10, 6))
sns.countplot(x='grade', hue='loan_status', data=df)
plt.title('Loan Status by Grade')
plt.show()
```

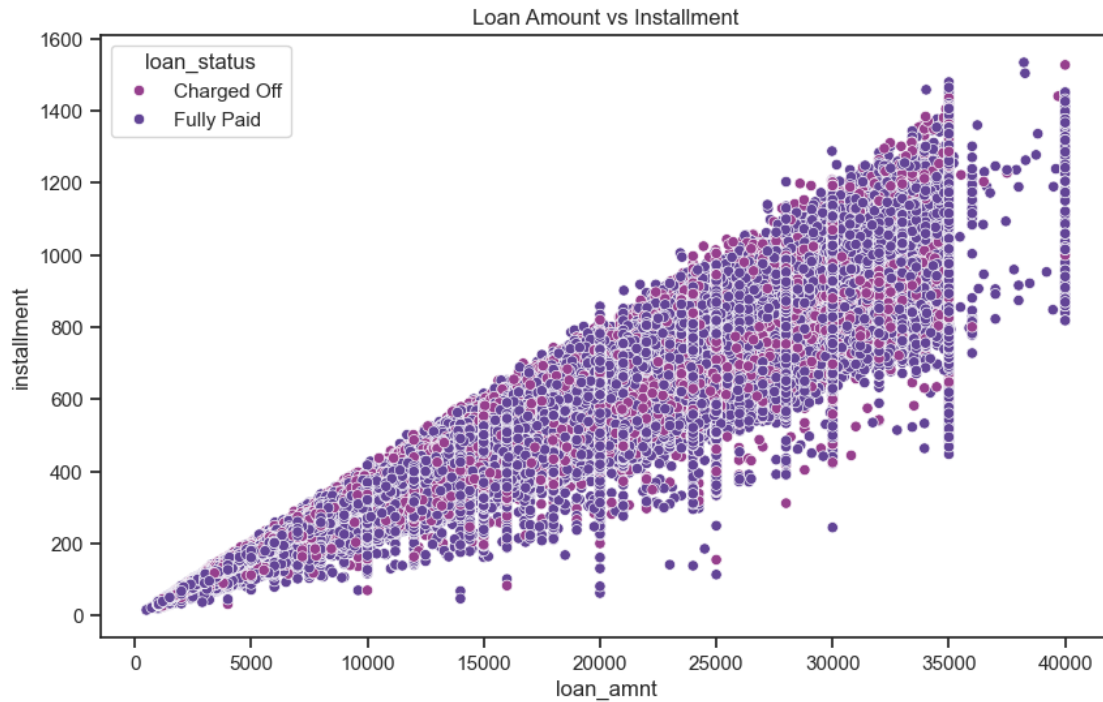


```
[28]: df_ = df[df['grade']=='A']
df_['loan_status'].value_counts(normalize=True) * 100
```

```
[28]: loan_status
Fully Paid      93.712122
Charged Off      6.287878
Name: proportion, dtype: float64
```

**People with grades 'A' are more likely to fully pay their loan. (T/F) - True, ~94%**  
 people with 'A' Grade paid their loan

```
[29]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='loan_amnt', y='installment', data=df, hue='loan_status')
plt.title('Loan Amount vs Installment')
plt.show()
```



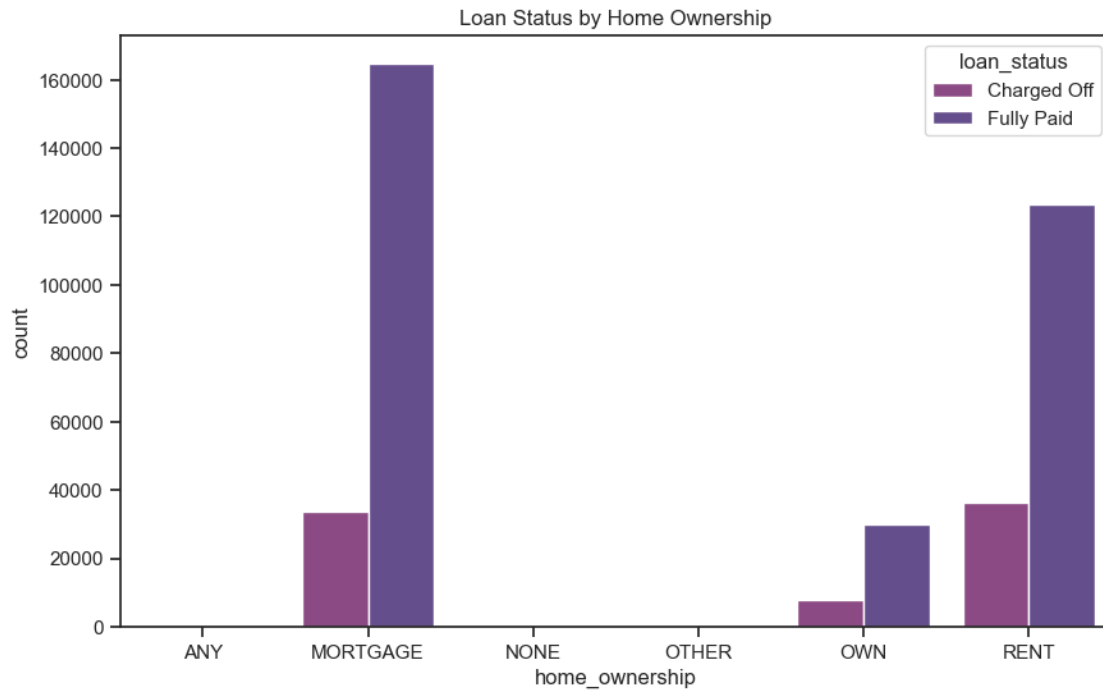
```
[30]: correlation = df['loan_amnt'].corr(df['installment'])
print(f"correlation between Loan Amount and Installment features: ", np.
      round(correlation, 2))
```

correlation between Loan Amount and Installment features: 0.95

Correlation between Loan Amount and Installment features: - 0.95

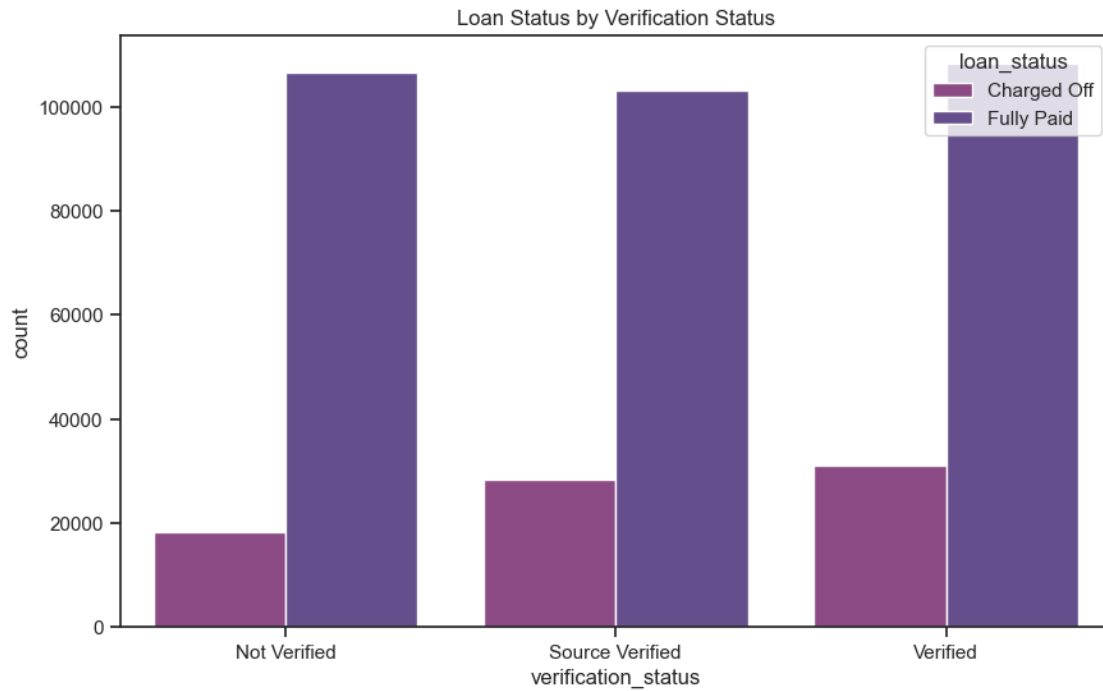
```
[31]: plt.figure(figsize=(10, 6))
sns.countplot(x='home_ownership', hue='loan_status', data=df)
plt.title('Loan Status by Home Ownership')
plt.show()
```



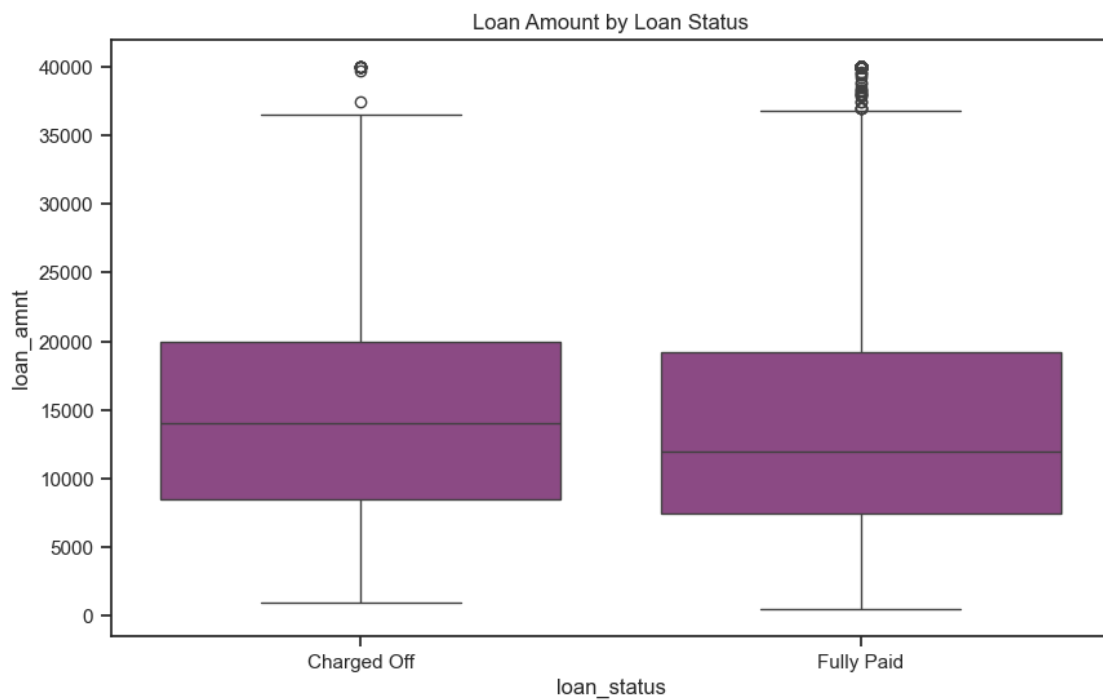


The majority of people have home ownership as - Mortgage

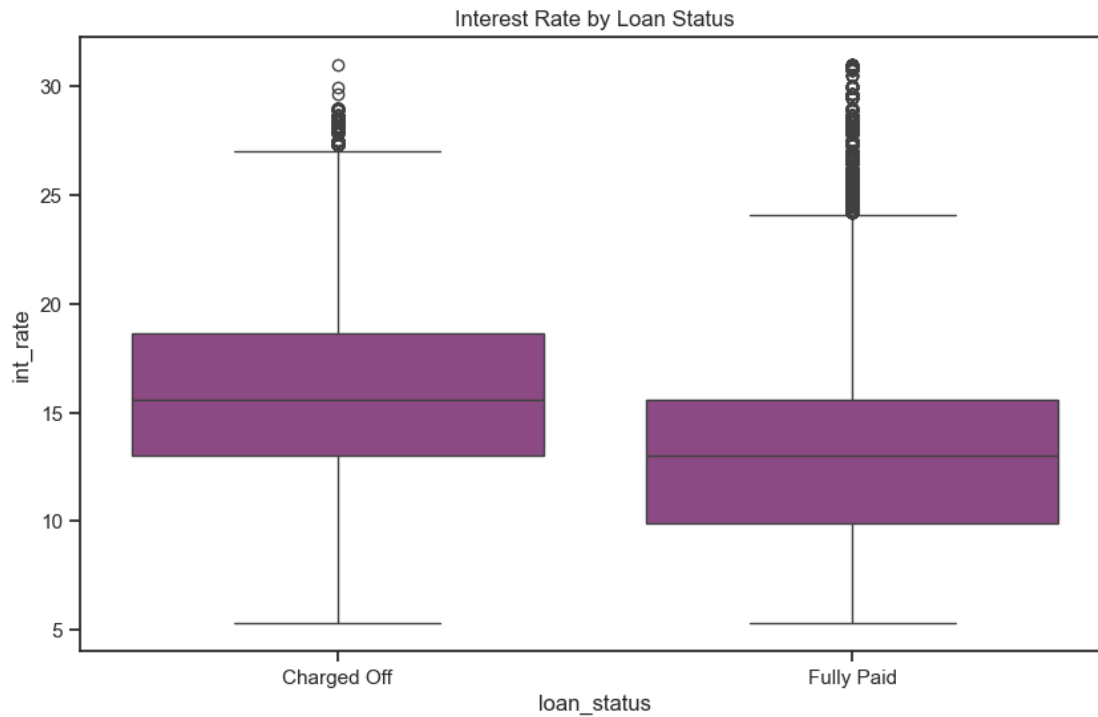
```
[32]: plt.figure(figsize=(10, 6))
sns.countplot(x='verification_status', hue='loan_status', data=df)
plt.title('Loan Status by Verification Status')
plt.show()
```




```
[33]: plt.figure(figsize=(10, 6))
sns.boxplot(x='loan_status', y='loan_amnt', data=df)
plt.title('Loan Amount by Loan Status')
plt.show()
```



```
[34]: plt.figure(figsize=(10, 6))
sns.boxplot(x='loan_status', y='int_rate', data=df)
plt.title('Interest Rate by Loan Status')
plt.show()
```



```
[35]: plt.figure(figsize=(10, 6))
sns.boxplot(x='loan_status', y='annual_inc', data=df)
plt.title('Annual Income by Loan Status')
plt.show()
```



```
[36]: numeric_df = df.select_dtypes(include=[np.number])
      corr_matrix = numeric_df.corr()

      plt.figure(figsize=(14, 10))
      sns.heatmap(corr_matrix, annot=True, cmap=palette, linewidths=0.5)
      plt.title('Correlation Matrix Heat Map')
      plt.show()
```



```
[37]: def detect_outliers_iqr(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return df[(df[column] < lower_bound) | (df[column] > upper_bound)]
```

### 0.1.9 Insights:

- Annual income columns has higher outliers
- More loans are availed for “Debt Consolidation”
- ~80% of loans opted for “Debt Consolidation” are Fully paid, ~20% are Charged off
- Second highest reason for loan for “Credit\_card”
- 0.95 correlation between Loan Amount and Installment features, Hence one can be removed during modelling

## 0.2 Data preparation for Modelling

```
[38]: df = df.drop(columns=['installment'])
```

```
[39]: X = df.drop(columns=['loan_status'])
# y = df['loan_status']
y = df['loan_status'].apply(lambda x: 1 if x == 'Fully Paid' else 0)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```
[40]: # preprocessing steps
numeric_features = ['loan_amnt', 'int_rate', 'annual_inc', 'dti', 'open_acc',
                    'revol_bal', 'revol_util', 'total_acc', 'pub_rec_flag', 'mort_acc_flag',
                    'pub_rec_bankruptcies_flag', 'zipcode']

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical_features = ['term', 'grade', 'sub_grade', 'home_ownership',
                        'verification_status', 'purpose', 'initial_list_status', 'application_type',
                        'state']

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(sparse_output=False, drop='first'))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

X_train = preprocessor.fit_transform(X_train)
```

```
[41]: model = LogisticRegression()
model.fit(X_train, y_train)
```

```
[41]: LogisticRegression()
```

```
[42]: X_test = preprocessor.transform(X_test)
y_pred = model.predict(X_test)
```

```
[43]: # Model Evaluation
from sklearn.metrics import accuracy_score, classification_report
```

```
print("Accuracy: ", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.8361235259955054
           precision    recall  f1-score   support

    0         0.69       0.31       0.43       15577
    1         0.85       0.97       0.90       63629

 accuracy
macro avg         0.77       0.64       0.66       79206
weighted avg         0.82       0.84       0.81       79206
```

```
[44]: from sklearn.metrics import precision_score, recall_score, f1_score, \
      ↪roc_auc_score

print("Precision: ", precision_score(y_test, y_pred))
print("Recall: ", recall_score(y_test, y_pred))
print("F1 Score: ", f1_score(y_test, y_pred))
print("ROC AUC Score: ", roc_auc_score(y_test, y_pred))
```

```
Precision: 0.8507208442394781
Recall: 0.9654088544531582
F1 Score: 0.9044435937453988
ROC AUC Score: 0.6367135432309445
```

**Thinking from a bank's perspective, which metric should our primary focus be on: -**  
 ROC AUC - Precision - Recall - F1 Score

**Bank's Perspective** From a bank's perspective, the primary focus should typically be on minimizing financial risk while maximizing opportunities for profit. This involves a delicate balance:

- *Minimizing False Negatives (Recall):*
  - Missing a defaulter can lead to financial losses, so high recall is important.
- *Minimizing False Positives (Precision):*
  - Incorrectly predicting a non-defaulter as a defaulter can lead to lost business opportunities, so high precision is also important.

### Recommended Primary Metric: *F1 Score*

- F1 Score is recommended as the primary metric because it balances precision and recall, which is essential for a bank's loan approval model. By focusing on the F1 Score, the bank can ensure that the model is effective at identifying defaulters while also minimizing the rejection of good loan applicants.

#### 0.2.1 Model Performance Metrics:

- Precision: 0.8507

- Recall: 0.9655
- F1 Score: 0.9045
- ROC AUC Score: 0.6368

## Understanding Precision and Recall

*Precision (0.8507)*: This indicates that 85.07% of the loans predicted as defaulters are actually defaulters. A precision of 0.8507 means that there is a 14.93% rate of false positives (non-defaulters incorrectly predicted as defaulters).

*Recall (0.9655)*: This indicates that 96.55% of actual defaulters are correctly identified by the model. A recall of 0.9655 means that there is a 3.45% rate of false negatives (defaulters incorrectly predicted as non-defaulters).

### How does the gap in precision and recall affect the bank?

The gap between precision and recall indicates that while your model is very good at identifying defaulters (high recall), it also incorrectly flags some non-defaulters as defaulters (lower precision).

```
[45]: # Feature Importance
importance = model.coef_[0]
numeric_features_list = preprocessor.transformers_[0][2]
categorical_features_list = preprocessor.named_transformers_['cat']['onehot'].
    ↪get_feature_names_out(categorical_features)
all_features = list(numeric_features_list) + list(categorical_features_list)
feature_importance = pd.DataFrame({'Feature': all_features, 'Importance':
    ↪importance})
feature_importance = feature_importance.sort_values(by='Importance',
    ↪ascending=False)
feature_importance
```

```
[45]:
```

	Feature	Importance
75	application_type_JOINT	1.754345
53	home_ownership_MORTGAGE	1.009141
56	home_ownership_OWN	0.859076
57	home_ownership_RENT	0.749367
55	home_ownership_OTHER	0.492414
72	purpose_wedding	0.425202
1	int_rate	0.355426
2	annual_inc	0.185882
54	home_ownership_NONE	0.135111
7	total_acc	0.118117
101	state_MN	0.108354
64	purpose_house	0.090313
23	sub_grade_B1	0.069224
5	revol_bal	0.060948
112	state_NY	0.057535
86	state_DE	0.054669
124	state_VT	0.054387
82	state_CA	0.047774



83	state_CO	0.047229
93	state_IN	0.046346
108	state_NH	0.038223
120	state_TN	0.031719
9	mort_acc_flag	0.029253
113	state_OH	0.027078
114	state_OK	0.023852
77	state_AK	0.023454
81	state_AZ	0.021813
123	state_VA	0.021500
115	state_OR	0.019864
89	state_HI	0.019571
85	state_DC	0.018238
99	state_ME	0.017817
100	state_MI	0.013318
106	state_ND	0.011249
96	state_LA	0.010096
105	state_NC	0.008725
118	state_SC	0.007135
73	initial_list_status_w	0.006770
102	state_MO	0.006451
87	state_FL	0.005309
121	state_TX	0.004956
10	pub_rec_bankruptcies_flag	0.004553
107	state_NE	0.000334
90	state_IA	-0.000048
109	state_NJ	-0.000378
117	state_RI	-0.000395
78	state_AL	-0.000963
92	state_IL	-0.001433
94	state_KS	-0.002406
125	state_WA	-0.009146
119	state_SD	-0.009293
98	state_MD	-0.009819
111	state_NV	-0.011719
110	state_NM	-0.012393
122	state_UT	-0.012841
126	state_WI	-0.013895
104	state_MT	-0.019889
8	pub_rec_flag	-0.023633
80	state_AR	-0.024443
95	state_KY	-0.027286
71	purpose_vacation	-0.028110
79	state_AP	-0.031839
91	state_ID	-0.035145
97	state_MA	-0.042179
76	state_AE	-0.043678

60	purpose_credit_card	-0.048063
84	state_CT	-0.050708
116	state_PA	-0.059509
103	state_MS	-0.063033
88	state_GA	-0.066932
127	state_WV	-0.067203
6	revol_util	-0.067940
28	sub_grade_C1	-0.070853
24	sub_grade_B2	-0.071530
0	loan_amnt	-0.072828
68	purpose_other	-0.076083
61	purpose_debt_consolidation	-0.083264
65	purpose_major_purchase	-0.089890
128	state_WY	-0.114570
59	verification_status_Verified	-0.116311
4	open_acc	-0.118873
63	purpose_home_improvement	-0.139503
58	verification_status_Source Verified	-0.161955
67	purpose_moving	-0.210400
66	purpose_medical	-0.213280
62	purpose_educational	-0.228592
25	sub_grade_B3	-0.254894
74	application_type_INDIVIDUAL	-0.259876
29	sub_grade_C2	-0.264055
69	purpose_renewable_energy	-0.271468
33	sub_grade_D1	-0.277595
19	sub_grade_A2	-0.331528
38	sub_grade_E1	-0.345315
43	sub_grade_F1	-0.368631
12	term_ 60 months	-0.375852
34	sub_grade_D2	-0.400465
26	sub_grade_B4	-0.406340
30	sub_grade_C3	-0.414956
35	sub_grade_D3	-0.455671
39	sub_grade_E2	-0.457699
3	dti	-0.474969
70	purpose_small_business	-0.488873
48	sub_grade_G1	-0.494554
44	sub_grade_F2	-0.506619
31	sub_grade_C4	-0.508829
49	sub_grade_G2	-0.526377
40	sub_grade_E3	-0.559836
36	sub_grade_D4	-0.572010
27	sub_grade_B5	-0.572913
51	sub_grade_G4	-0.592506
20	sub_grade_A3	-0.597591
32	sub_grade_C5	-0.611243

45	sub_grade_F3	-0.653658
37	sub_grade_D5	-0.656049
41	sub_grade_E4	-0.660050
21	sub_grade_A4	-0.702664
46	sub_grade_F4	-0.721955
42	sub_grade_E5	-0.738930
52	sub_grade_G5	-0.763358
11	pincode	-0.854771
47	sub_grade_F5	-0.882899
22	sub_grade_A5	-0.957883
50	sub_grade_G3	-1.017774
13	grade_B	-1.236452
14	grade_C	-1.869935
15	grade_D	-2.361790
16	grade_E	-2.761831
17	grade_F	-3.133763
18	grade_G	-3.394571

Which were the features that heavily affected the outcome? \* application\_type \* home\_ownership \* purpose \* int\_rate \* annual\_inc \* total\_acc \* grade \* sub\_grade \* state

Will the results be affected by geographical location? - Yes. States features are listed in the top features

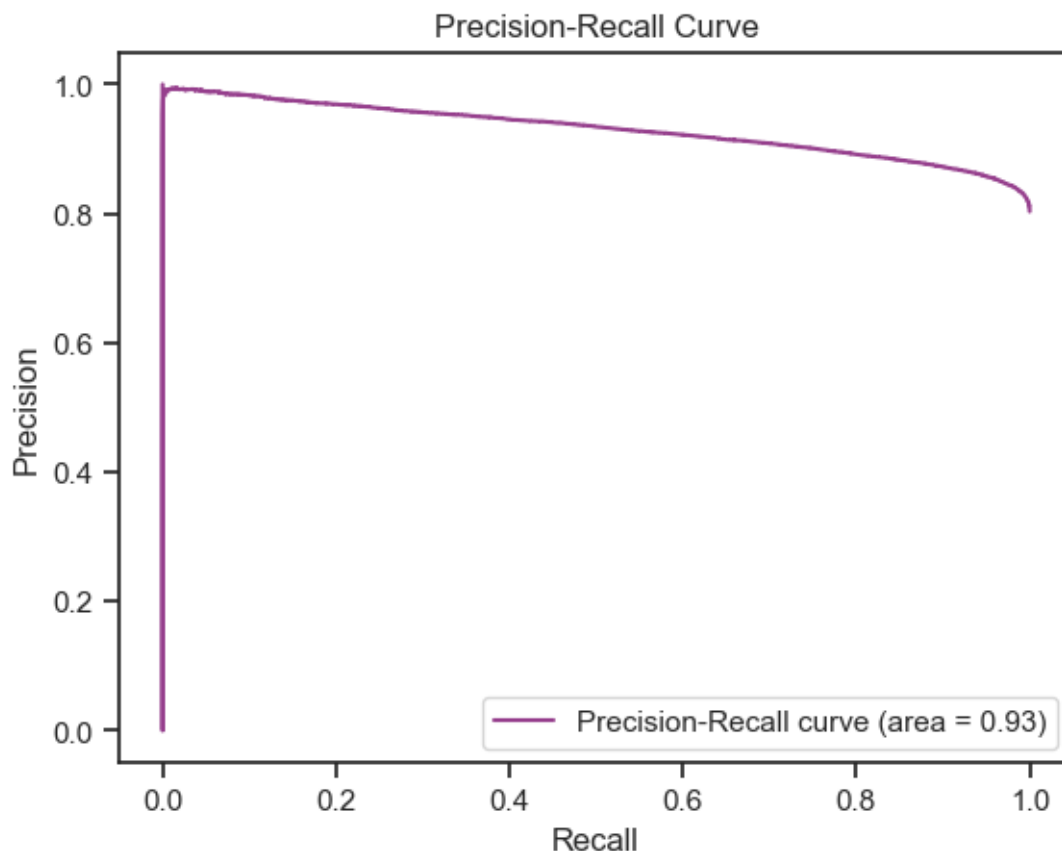
```
[46]: import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_curve, auc

# Predict probabilities
y_probs = model.predict_proba(X_test)[:, 1]

# Calculate precision-recall curve
precision, recall, _ = precision_recall_curve(y_test, y_probs)

# Calculate AUC
pr_auc = auc(recall, precision)

# Plot the precision-recall curve
plt.figure()
plt.plot(recall, precision, label=f'Precision-Recall curve (area = {pr_auc:.2f})')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='best')
plt.show()
```



### 0.2.2 Actionable Insights

#### High Recall, Lower Precision:

*Insight:* Your model is highly effective at identifying defaulters (high recall) but has a moderate rate of false positives (lower precision).

*Impact:* While the bank minimizes the risk of defaults, it may also be rejecting a significant number of creditworthy applicants, leading to lost business opportunities.

#### Geographical Variability:

*Insight:* Geographical differences can significantly impact loan default rates due to varying economic conditions, cost of living, and access to financial services.

*Impact:* Ignoring geographical factors may lead to suboptimal loan approval decisions and higher default rates in certain regions.

### 0.2.3 Recommendations

#### Segmented Modeling:

Develop separate models for different geographical regions or clusters with similar characteristics. This allows for more tailored predictions and better handling of regional variations.

**Post-Processing Rules and Manual Review:**

Implement post-processing rules or manual reviews for borderline cases to reduce false positives without significantly impacting recall.

---

[ ]: